

# Target-Guided AutoEncoder

---

Niclas Wölner-Hanssen, VT 2023, STAN47 – Deep Learning

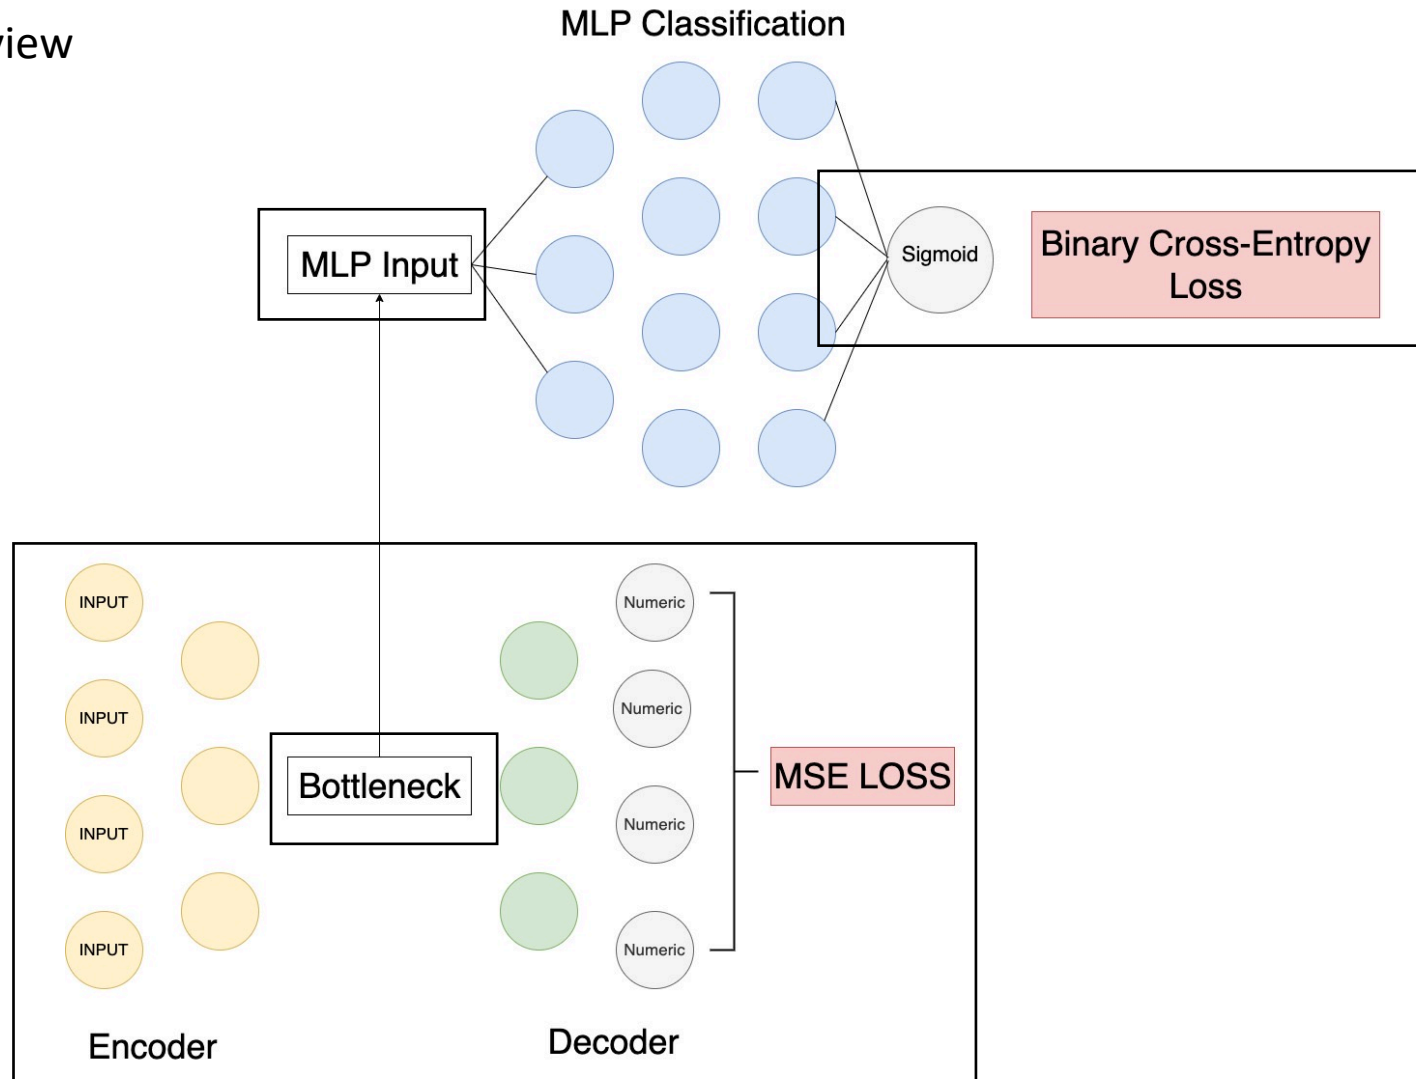


# Targeted-Guided AutoEncoder

- Feature extraction/ dimension reduction technique that is guided to construct latent variables that are informative.
- Often times the indirect purpose of dimension reduction techniques such as PCA and AE is to reduce the number of dimensions to make better predictions (overcome multicollinearity, overfitting etc.)
- However, these transformed new features, the latent variables, are not necessarily predictive.
- No information about the target, the thing we try to predict, is introduced in the transformation of the features into the latent variables when using PCA or Simple AutoEncoder.

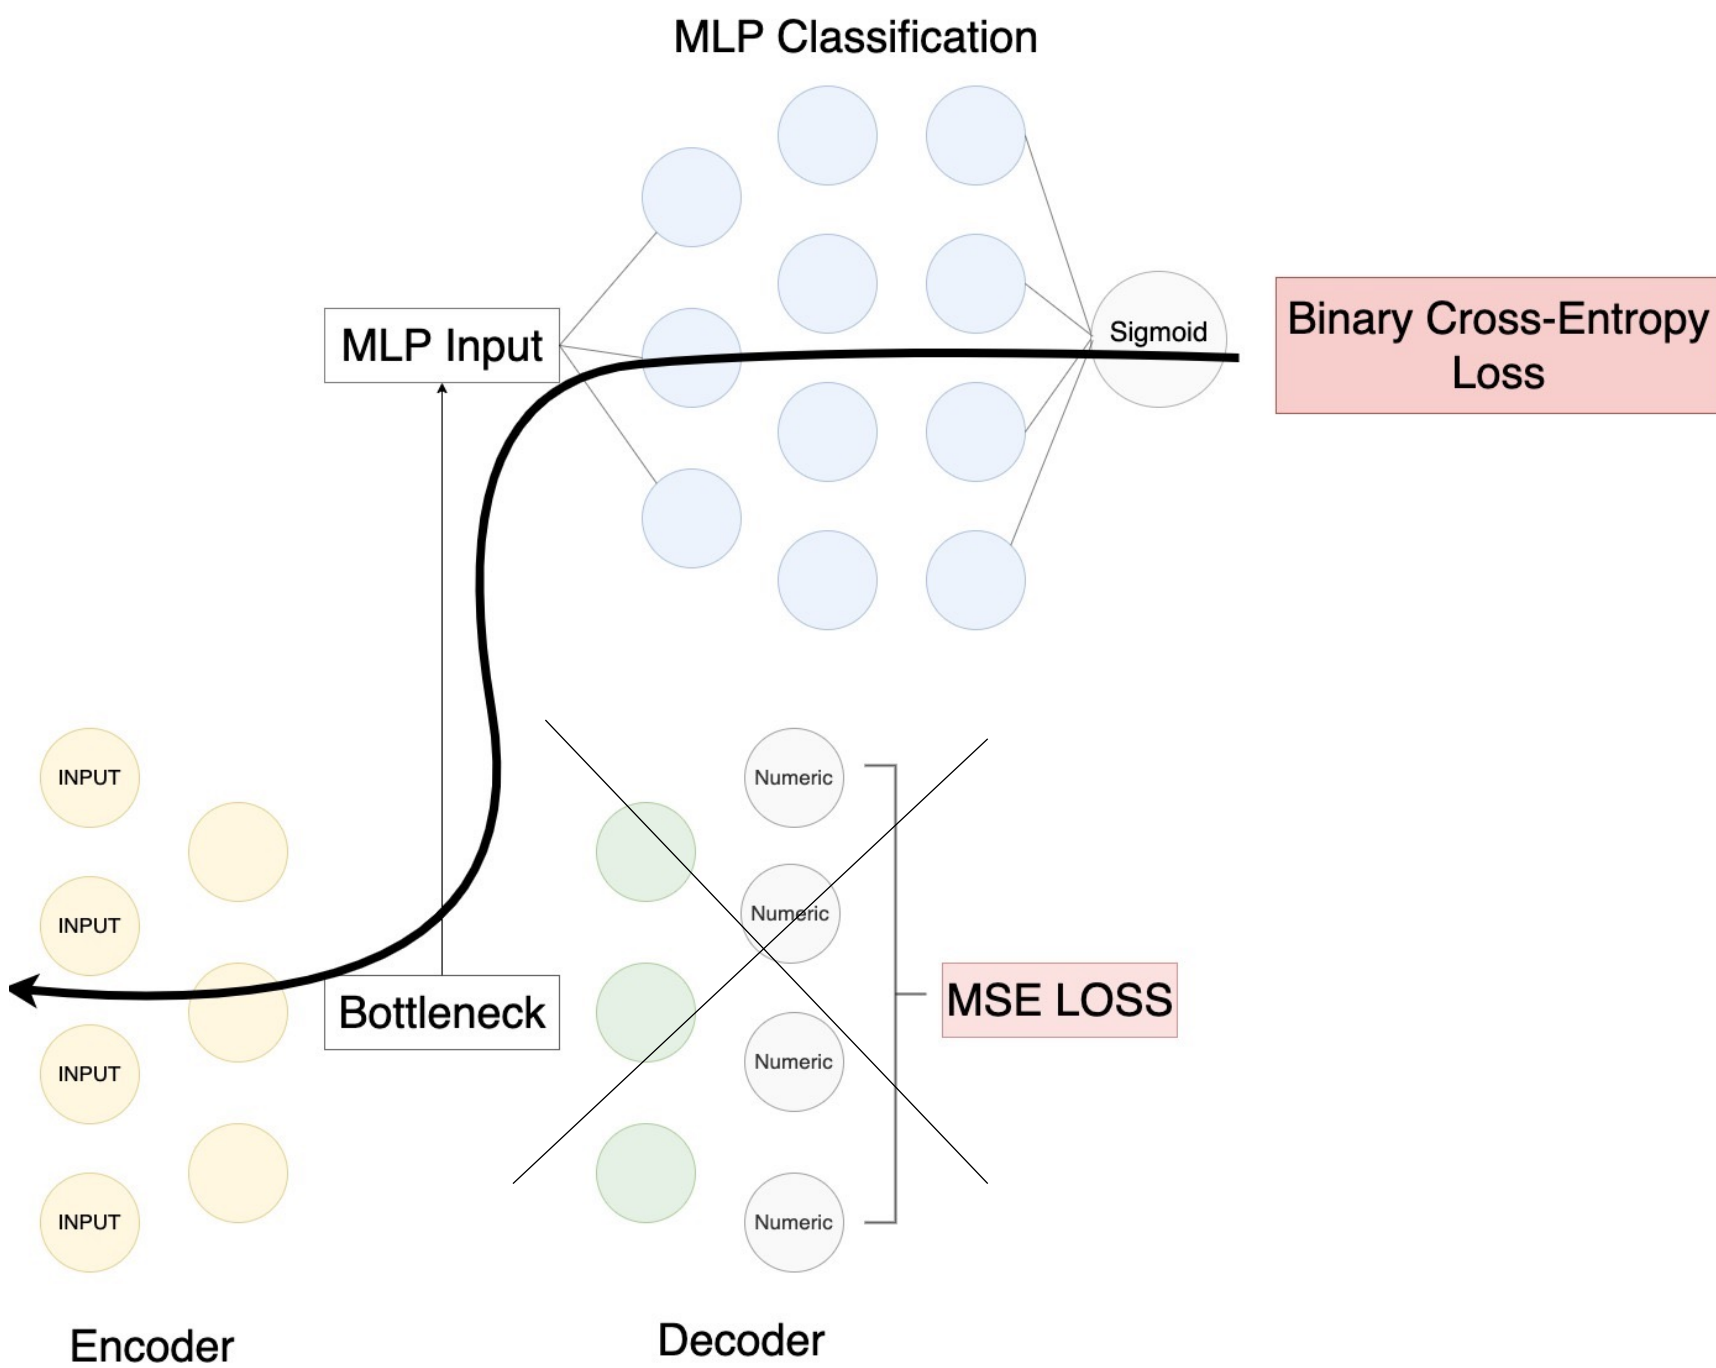
# Targeted-Guided AutoEncoder


## Model Overview



# Target-Guided AutoEncoder

- By including a separate MLP branch in the overall model architecture, the target information (i.e  $y_{\text{train}}$ ) is introduced to the shared bottleneck layer of the autoencoder.
- During training, the gradients of the MLP loss with respect to the weights of the MLP are propagated backwards through the shared bottleneck layer of the autoencoder, allowing the model to update the bottleneck representation in a way that is informed by the classification task.
- Bottomline is that the inclusion of the MLP branch enables the model to learn the bottleneck representation that is not only optimized for reconstruction of the input data but also takes the target information into account.





MLP Input

MLP Classification

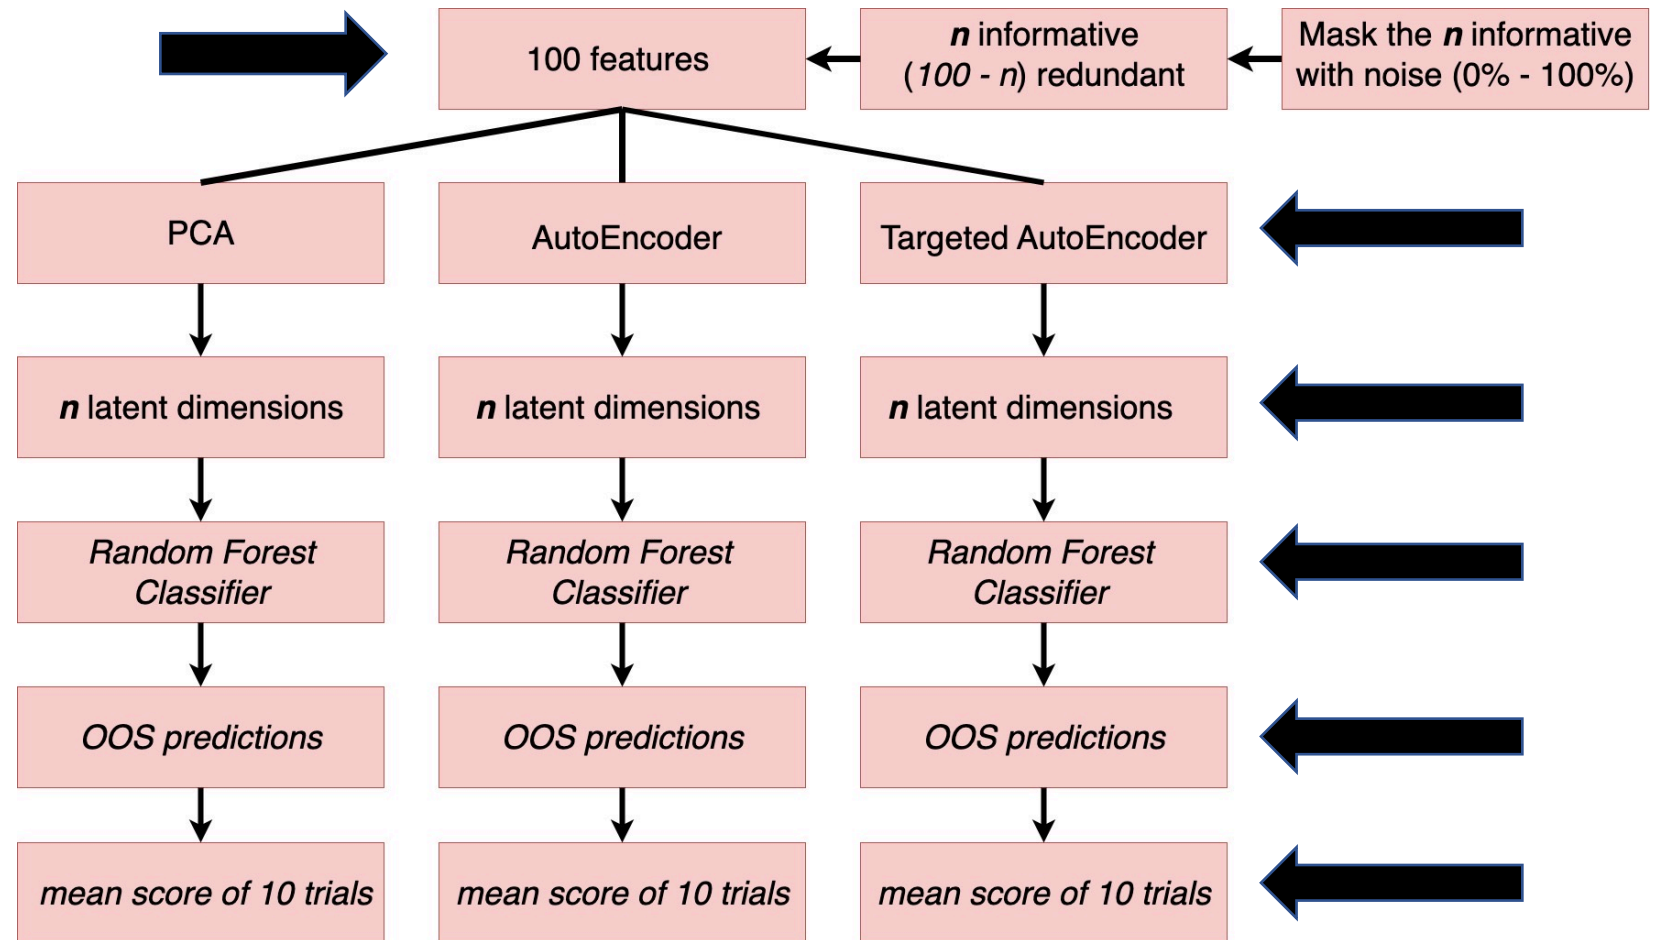
Sigmoid

Binary Cross-Entropy Loss

Bottleneck

MSE LOSS


# Testing



\*RF:  $n\_estimators = 500$ ,  $max\_depth = 5$ ,  $max\_samples = 100$




## Testing



- 100 features (IID Std Normal, 10K samples)
- The sum of the first 3 features makes up the binary target:
  - $y > 0 = 1$
  - $y < 0 = 0$
- The other features are thus just noise.
- In the plot the 3 informative features are after the construction of the target masked with noise. Going from complete noise (100%) to no-noise (0%)(x-axis).

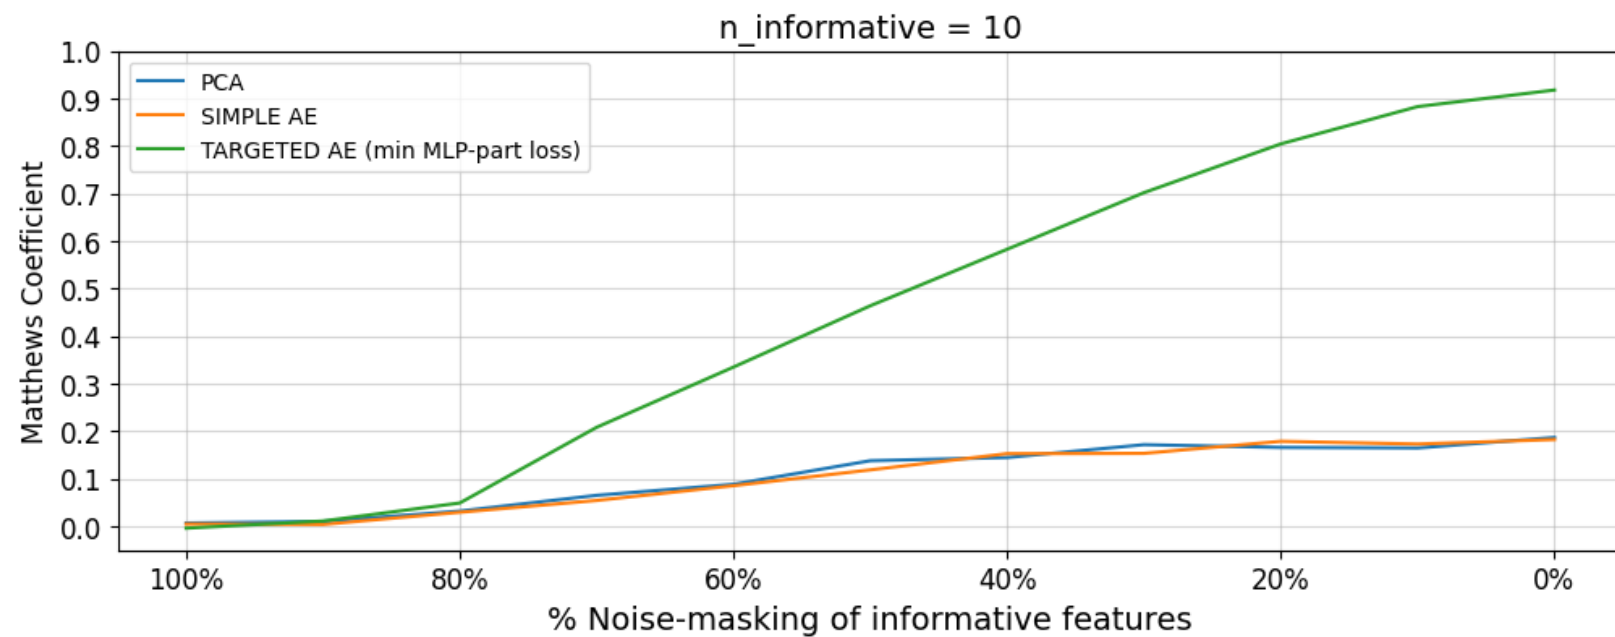



## Testing

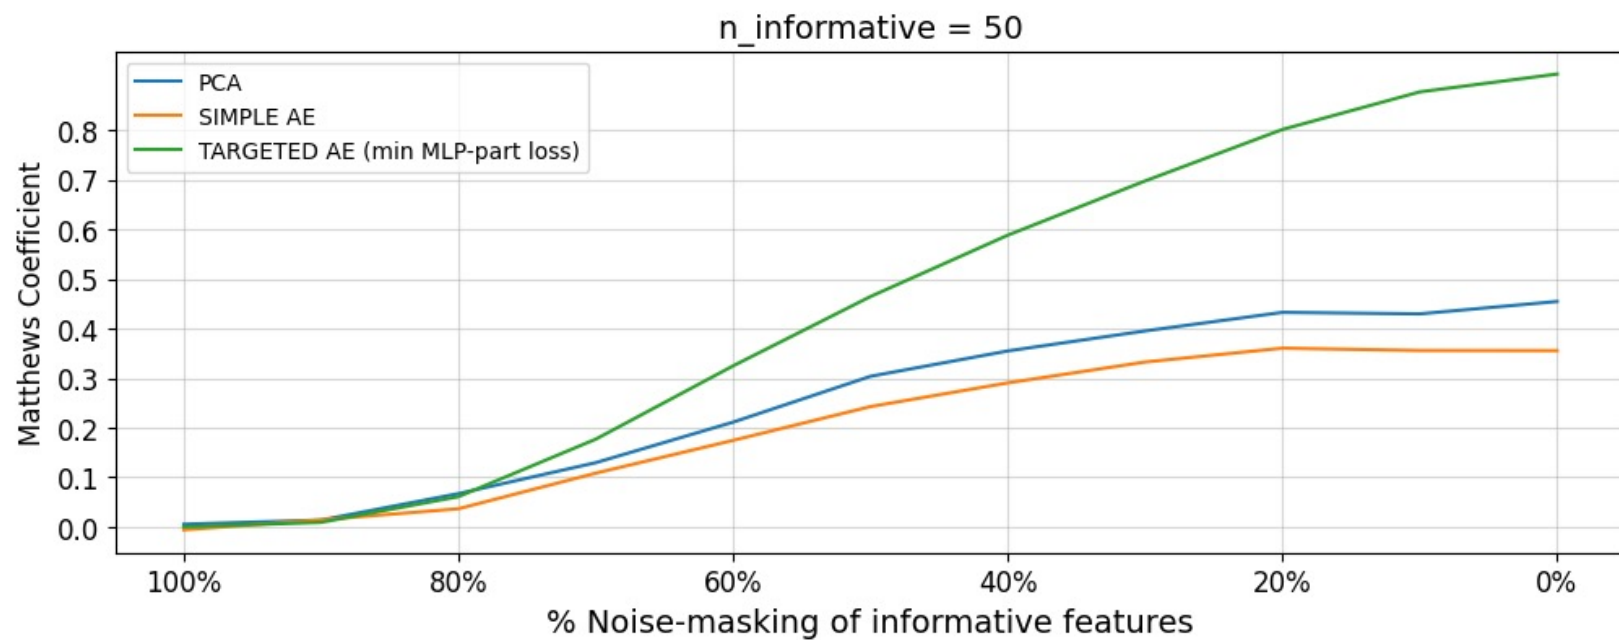



- Matthews correlation coefficient

# Testing




# Testing



# Target-Guided AutoEncoder

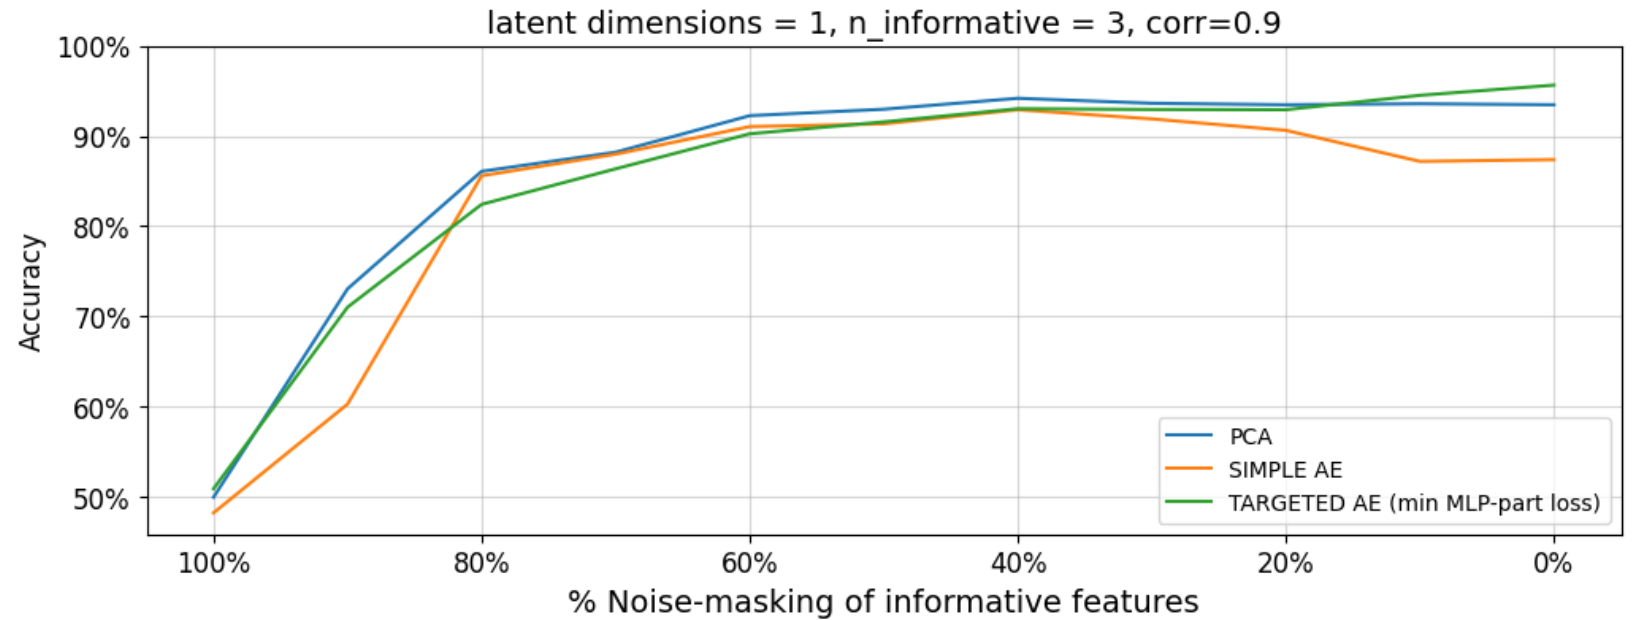
- So far, the samples has been IID. However, PCA is based on the covariance/correlation matrix and generally the higher the correlation between the features the more information (i.e. variance) of these it can store in fewer dimensions.
- For this reason, lets introduce some correlation between the features.

## Testing




- The target is constructed as the sign of the sum of the 3 informative features, the other 97 features are "redundant".
- However, the correlation between all features is approx 0.2 when noise = 0%.
- PCA: First principal component
- AE & Targeted AE: 1 latent dimension

## Testing



- The target is constructed as the sign of the sum of the 3 informative features, the other 97 features are "redundant".
- However, the correlation between all features is approx 0.9 when noise = 0%.
- PCA: First principal component
- AE & Targeted AE: 1 latent dimension


## Testing



- Within the 3 informative features the correlation is 0.9 when noise = 0%.
- Within the 97 redundant features the correlation is 0.9 when noise = 0%.
- Between these groups the correlation is approx 0.
- PCA: First principal component
- AE & Targeted AE: 1 latent dimension



## Testing



- Within the 50 informative features the correlation is 0.9 when noise = 0%.
- Within the 50 redundant features the correlation is 0.9 when noise = 0%.
- Between these groups the correlation is approx 0.
- PCA: First principal component
- AE & Targeted AE: 1 latent dimension

# Conclusion:



Testing

- The Target-Guided AutoEncoder tend to at least never underperform PCA and simple AE, and most often it outperforms in terms of informative quality in constructing latent variables.
- The results also shows that PCA and Simple AutoEncoder seems to from time to time miss the informative information in the features, whereas the Target-Guided AutoEncoder finds these with realtive ease most of the time.

# Code

```
In [94]: 1 # input regularization
2 input_x = Input(shape=(x_inputs,))
3 input_x = GaussianNoise(0.03)(input_x)
4 input_x = BatchNormalization()(input_x)
5 # encoder level 1
6 e = Dense(x_inputs, activation='tanh', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01))(input_x) #layer 1
7 e = BatchNormalization()(e)
8 e = Dropout(dropout_rate)(e)
9 # encoder level 2
10 e = Dense(int(x_inputs/2), activation='tanh')(e) #layer 2
11 e = BatchNormalization()(e)
12 e = Dropout(dropout_rate)(e)
13 # bottleneck
14 bottleneck = Dense(latent_dimensions)(e) #Bottleneck
15
16 # decoder, level 1
17 d = Dense(int(x_inputs/2), activation='tanh')(bottleneck) #layer 1
18 d = BatchNormalization()(d)
19 d = Dropout(dropout_rate)(d)
20 # decoder level 2
21 d = Dense(x_inputs, activation='tanh')(d) #layer 2
22 d = BatchNormalization()(d)
23 d = Dropout(dropout_rate)(d)
24 # output layer
25 decoder_output = Dense(x_inputs, activation='linear', name = 'decoder_output')(d)
26
27 merged = concatenate([bottleneck])
28
29 #####MLP Classification
30
31 mlp = Dense(merged.shape[1], activation='tanh')(merged) #Input Layer MLP
32 mlp = BatchNormalization()(mlp)
33 mlp = Dropout(dropout_rate)(mlp)
34
35 mlp = Dense(x_inputs*3, activation='tanh')(mlp) #layer 1
36 mlp = BatchNormalization()(mlp)
37 mlp = Dropout(dropout_rate)(mlp)
38
39 mlp = Dense(x_inputs*3, activation='tanh')(mlp) #layer 2
40 mlp = BatchNormalization()(mlp)
41 mlp = Dropout(dropout_rate)(mlp)
42
43 mlp = Dense(x_inputs*2, activation='tanh')(mlp) #layer 3
44 mlp = BatchNormalization()(mlp)
45 mlp = Dropout(dropout_rate)(mlp)
46
47 mlp_output = Dense(1, activation = 'sigmoid', name = 'mlp_output')(mlp) #Output Layer MLP
48
49 # define autoencoder model
50 model = Model(inputs= input_x, outputs=[decoder_output, mlp_output])
51 encoder = Model(inputs= input_x, outputs=bottleneck)
52 decoder = Model(inputs= bottleneck, outputs=decoder_output)
53 MLP = Model(inputs= merged, outputs=mlp_output)
54
55 # compile model
56 model.compile(optimizer = tf.keras.optimizers.Adam(learning_rate = 1e-3),
57               loss= {'decoder_output': 'mse', 'mlp_output': 'binary_crossentropy'},
58               metrics={'decoder_output': 'mse', 'mlp_output': AUC()},
59               sample_weight_mode={'decoder_output': None, 'mlp_output': 'temporal'},
60               weighted_metrics={'mlp_output': 'binary_crossentropy'})
61
```

Encoder

Decoder

MLP



## Code



```
es = EarlyStopping(monitor = 'val_mlp_output_loss', min_delta = 1e-4, patience = 20,  
                  baseline = None, restore_best_weights = True, verbose = 0)  
  
history = model.fit(  
    x = X_train,  
    y = {'decoder_output': X_train, 'mlp_output': y_train},  
    epochs=200,  
    batch_size=100,  
    verbose=0,  
    shuffle=True,  
    validation_split = 0.2,  
    callbacks = [ckp, es])
```



Thank you!